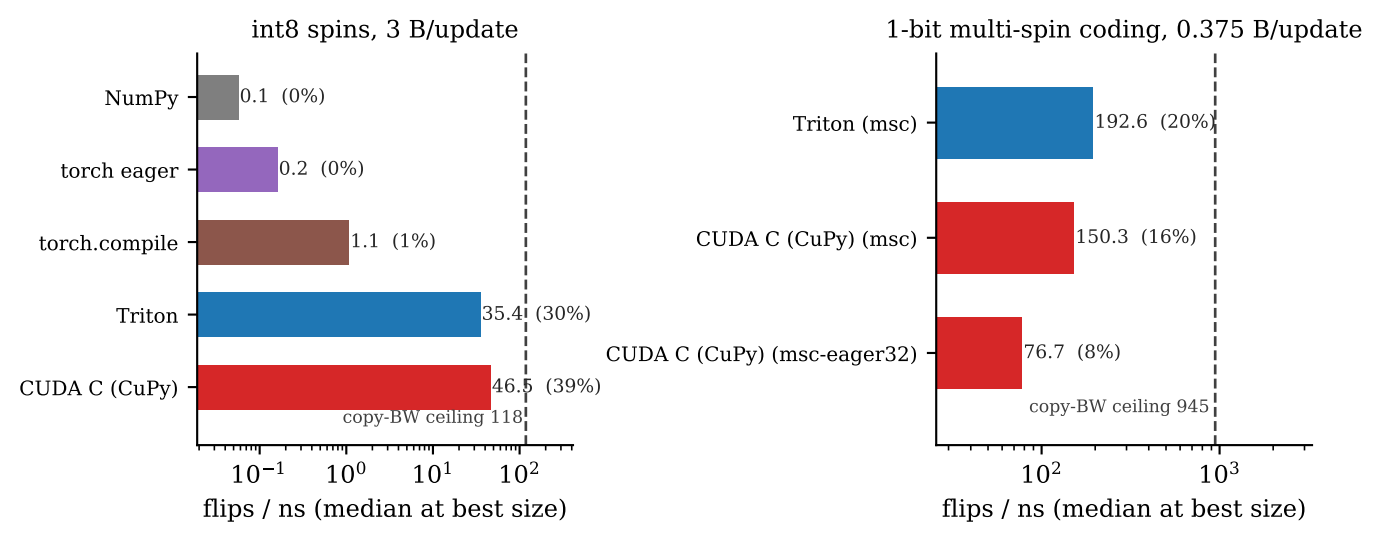
int8 spins, 3 B/update
NumPy
0.1 (0%)
torch eager
0.2 (0%)
torch.compile
1.1 (1%)
Triton
35.4 (30%)
CUDA C (CuPy)
46.5 (39%)
copy-BW ceiling 118
10^{-1}
10^{0}
10^{1}
10^{2}
flips / ns (median at best size)
1-bit multi-spin coding, 0.375 B/update
Triton (msc)
192.6 (20%)
CUDA C (CuPy) (msc)
150.3 (16%)
CUDA C (CuPy) (msc-eager32)
76.7 (8%)
copy-BW ceiling 945
10^{2}
10^{3}
flips / ns (median at best size)