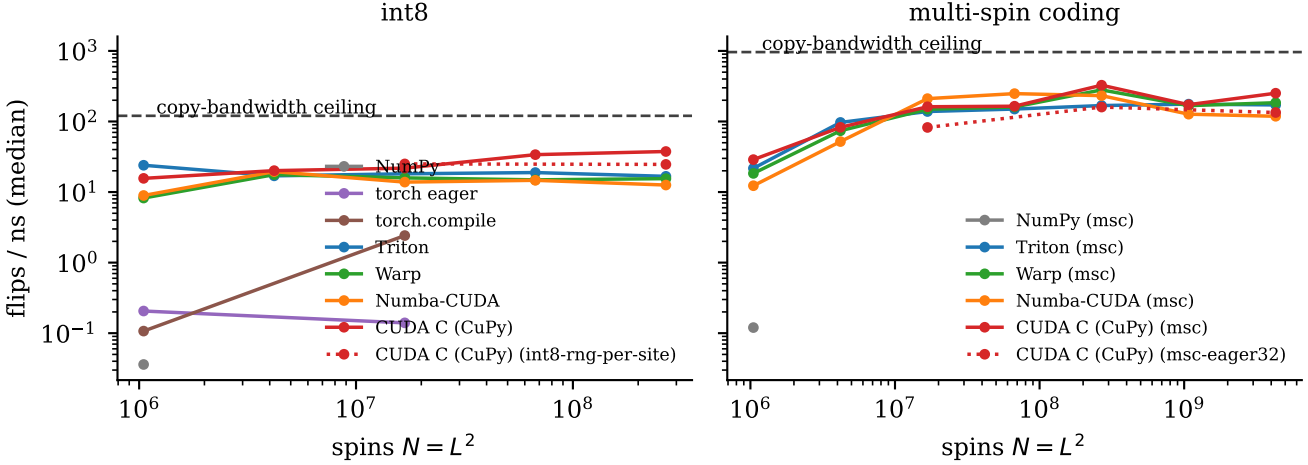
int8
multi-spin coding
flips / ns (median)
copy-bandwidth ceiling
copy-bandwidth ceiling
NumPy
torch eager
torch.compile
Triton
Warp
Numba-CUDA
CUDA C (CuPy)
CUDA C (CuPy) (int8-rng-per-site)
NumPy (msc)
Triton (msc)
Warp (msc)
Numba-CUDA (msc)
CUDA C (CuPy) (msc)
CUDA C (CuPy) (msc-eager32)
spins $N = L^2$
spins $N = L^2$